



Creating an EKS cluster in AWS



Terence Wong

February 9, 2022 • 3 mins

In this post, you learn how to set up an Elastic Kubernetes Service (EKS) cluster in Amazon Web Services (AWS).

EKS is a managed container service to run and scale Kubernetes in the cloud or on-premises. Kubernetes provides a scalable, distributed way to manage workloads. It does this by containerizing applications. Containers ensure replicability across different environments and cloud infrastructures.

The clusters you create in this post will be used in later posts in our [Continuous Integration series](#), to set up web applications and as part of workflows. We'll add links to the relevant posts as they become available:

- [Building a Docker image in Jenkinsfile and publishing to ECR](#)
- [Deploying to Amazon EKS with Docker and Jenkins](#)

Prerequisites

To follow along with this post, you need:

- An AWS account
- A terminal with [kubect](#), [eksctl](#), and [aws-iam-authenticator](#) installed
- IAM permissions

There are two ways to set up a cluster in EKS:

- The command-line interface (CLI)
- The console

Before this though, you need to set up some access keys and user accounts.

Preliminary setup

In AWS, you need to configure access policies. These policies determine what kind of user can access the cluster.

Typically, it's useful to follow the principle of least privilege. This means you give users the minimum set of privileges required to carry out their role. This supports the security of the cloud infrastructure by not granting users more permissions than they need.

Follow these steps to set up your access keys and user accounts:

- Go to **AWS Console**, then **IAM**, then **Users**, and **Add Users**
- Give the user a name, and tick **Access Key - Programmatic access**
- Click **Next**

Add user

1

2

3

4

5

Set user details

You can add multiple users at once with the same access type and permissions. [Learn more](#)

User name*

test.user

+

Add another user

Select AWS access type

Select how these users will primarily access AWS. If you choose only programmatic access, it does NOT prevent users from accessing the console using an assumed role. Access keys and autogenerated passwords are provided in the last step. [Learn more](#)

Select AWS credential type*

☒

Access key - Programmatic access
Enables an **access key ID** and **secret access key** for the AWS API, CLI, SDK, and other development tools.

☐

Password - AWS Management Console access
Enables a **password** that allows users to sign-in to the AWS Management Console.

- Select **Attach existing policies directly**, then **Create policy**

The IAM policy allows you to create an EKS cluster from the command-line. [The actions in this policy are the minimum policies required by eksctl.](#)

- Click **Next** and give the policy a name
- Click **Add Policy**

Amazon then shows you the **Access Key ID** and **Secret Access Key**. Download this file for reference later.

Command-line interface

Log in to the AWS command-line using `aws login`.

Run `aws configure`.

Enter your Access Key ID and Secret Access Key earlier. Set the zone to `us-east-2` and accept the defaults.

Now you can create your cluster.

```
eksctl create cluster \  
--name my-cluster \  
--region us-east-2 \  
--fargate
```

AWS Fargate lets you run containers without managing servers or clusters of Amazon EC2 instances. The profile provides a simple way to spin up a test cluster. Test the cluster configuration:

```
kubectl get svc
```

Output

<https://octopus.com/blog/eks-cluster-aws>

2/5

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.100.0.1	<none>	443/TCP	25h

Run the following command to view your cluster nodes:

```
kubectl get nodes -o wide
```

Output

NAME	STATUS	ROLES	AGE	VERSION
fargate-ip-192-168-129-76.us-east-2.compute.internal	Ready	<none>	25h	v1.20.7-eks-1353
fargate-ip-192-168-165-146.us-east-2.compute.internal	Ready	<none>	25h	v1.20.7-eks-1353

View the workloads that are running on your cluster:

```
kubectl get pods --all-namespaces -o wide
```

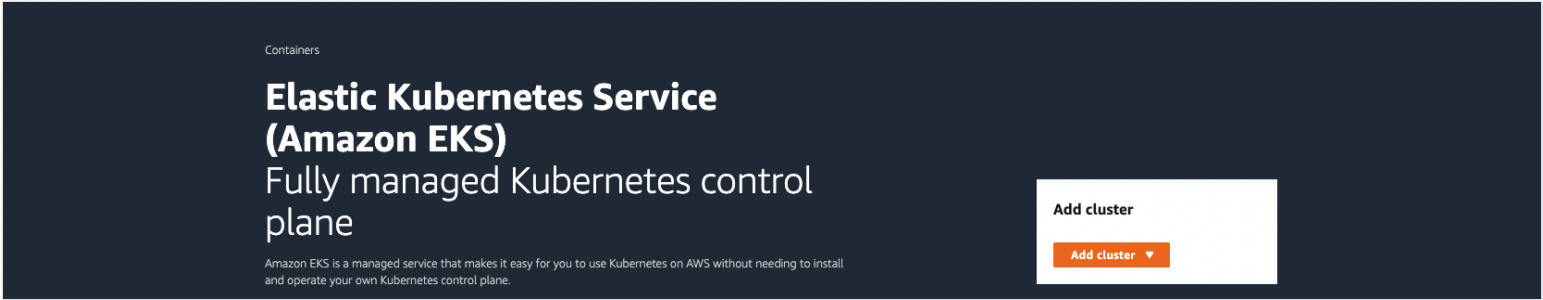
Output

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE	IP	NO
kube-system	coredns-85f9f6cd8b-2n8wr	1/1	Running	0	25h	192.168.129.76	fa
kube-system	coredns-85f9f6cd8b-c4jfk	1/1	Running	0	25h	192.168.165.146	fa

Console interface

You can also create a cluster from the AWS console.

Go to **EKS**, then **Add Cluster**, then **Create**.



- **Name:** give your cluster a name
- **Kubernetes Version:** select the latest Kubernetes version
- **Cluster Service Role:** re-use the service role created from the CLI cluster

Accept all other defaults to create your cluster.

Checking clusters

Check the state of your two clusters by going to **EKS**, then **Clusters**.

	Cluster name	Status	Kubernetes version	Provider
<input type="radio"/>	my-cluster-cli	Active	1.20 Update now	EKS
<input type="radio"/>	my-cluster-console	Active	1.21	EKS

Now that your clusters are live, you can perform operations on them, such as deploying an application or configuring resources. For this example, we'll delete them.

Deleting the clusters

You can delete a cluster using the CLI by running this command. Replace the cluster name and region with your values.

```
eksctl delete cluster --name my-cluster --region us-east-2
```

You can delete a cluster using the console by selecting the cluster, clicking **Delete**, and typing the name of the cluster to delete.

You can only delete a resource the same way you created it. This means clusters created through the CLI can only be deleted through the CLI. Clusters created through the console can only be deleted through the console.

Clusters (2) Info

⌂

Delete

Add cluster ▾

🔍

Filter cluster by name, status, kubernetes version, or provider

< 1 >

	Cluster name	Status	Kubernetes version	Provider
<input type="radio"/>	my-cluster-cli	✔ Active	1.20 Update now	EKS
<input checked="" type="radio"/>	my-cluster-console	✔ Active	1.21	EKS

Conclusion

In this post, you set up IAM permissions in AWS. You used the CLI and console to create, inspect, and delete an EKS cluster.

EKS on AWS allows you to provision Kubernetes services in the cloud to deploy and scale workloads.

The clusters you created can be used to set up web applications and as part of workflows in upcoming posts. We'll add links to the relevant posts as they become available:

- [Building a Docker image in Jenkinsfile and publishing to ECR](#)
- [Deploying to Amazon EKS with Docker and Jenkins](#)

Read the rest of our [Continuous Integration series](#).

Happy deployments!

Tagged with:

DevOps


CI Series

Continuous Integration

AWS

Kubernetes

Related posts



Common mistakes in DevOps metrics



Steve Fenton
December 7, 2022 • 6 mins

Comparing Lean, Agile, and Continuous Delivery



Steve Fenton
December 5, 2022 • 6 mins



DevOps

Best practices for CI/CD



Terence Wong
November 30, 2022 • 5 mins

Newsletter

Logged in as Terence Wong (terence.wong@octopus.com)

Join ~48,000 DevOps professionals and sign up for the latest Octopus news, events, and opinions. No spam. Unsubscribe at any time.

Subscribe

Your privacy is important to us. Read more in our [Privacy Policy](#).

